

Documentation for Algorithms for Massive Datasets project: Turkish lira recognizer

Caccaro Sebastiano
Cavagnino Matteo
A.A.2019/2020

We declare that this material, which We now submit for assessment, is entirely our own work and has not been taken from the work of others, save and to the extent that such work has been cited and acknowledged within the text of our work. We understand that plagiarism, collusion, and copying are grave and serious offences in the university and accept the penalties that would be imposed should we engage in plagiarism, collusion or copying. This assignment, or any part of it, has not been previously submitted by us or any other person for assessment on this or any other course of study.

Contents

1	Introduction	1
1.1	References and sources	1
1.2	Links	1
2	Dataset	2
2.1	Images	2
2.2	Preprocessing	3
2.2.1	Image Scaling	3
2.2.2	Dataset loading and batch splitting	3
3	Experiments and results	4
3.1	Model summary	4
3.2	Models	5
3.2.1	Baseline Model	5
3.2.2	Convolution Model	6
3.2.3	Convolution and Pooling Model	8
3.2.4	Convolution and Pooling Model with Dropout	9
3.2.5	Batch Normalization Model	11
3.3	Models comparison	12
4	Conclusions	14
	References	15

1 Introduction

The goal of this project is to build and train a classifier for Turkish Liras based on the *Turkish Lira Banknote Dataset* (section 2).

The classifier is developed using Convolutional Neural Networks (from now on CNN). The CNN is coded using python3 and using the Tensorflow and Keras libraries.

1.1 References and sources

In order to develop the project some other sources other than the ones given during the course were consulted: all the links are cited in the *Reference* section of this document.

The sources include also some links from the *Tensorflow official documentation* from which most of the code in the project has been adapted.

1.2 Links

All of the code from the project is available at the following link:

https://github.com/sebacaccaro/Progetto_AFMD

The code is available both in a python file and in a notebook which is meant to be run on Google Colab.

Note that in both cases a Kaggle API key is needed in order download the dataset and run the code.

2 Dataset

The dataset used for the project is the *Turkish Lira Banknote Dataset*, available at the following link:

<https://www.kaggle.com/baltacifatih/turkish-lira-banknote-dataset>

The dataset is composed by 6000 images of Turkish Lira Banknotes, organized as follows:

- 1000 pictures of 5₺ banknotes
- 1000 pictures of 10₺ banknotes
- 1000 pictures of 20₺ banknotes
- 1000 pictures of 50₺ banknotes
- 1000 pictures of 100₺ banknotes
- 1000 pictures of 200₺ banknotes

Images are already splitted in *train* and *validation*. The training set contains 925 images for each banknote, leading to 92.5% of the images (5550) begin used for training and the remaining 7.5% (450) being used for validation.

2.1 Images

All the pictures in the dataset consist of picture of banknotes taken with different perspectives and conditions: some are just laying on different surfaces, some are held in hand, blurred, folded or partially out of frame, ecc.




Figure 1: Example of images of 20₺ banknotes

The author of dataset already performed brightness adjustment, noise reduction and image flipping when necessary.
All images are in PNG format with a 1280 x 720 resolution.

2.2 Preprocessing

Taking into consideration what has already been said about the images, it is clear that there is no need to apply other forms of corrections. Therefore, the only steps necessary prior to the train phase are:

- Image scaling and decoding
- Dataset loading and batch splitting

2.2.1 Image Scaling

All the images come in a 1280x720 resolution. Working with such high resolution it's not recommended when training a CNN as it would take a big toll on the training speed due to the increasing number of parameters and RAM usage. For this reason all the images are reduced by a scale factor of 5, adjusting their resolution to 140x256.

Every image is also decoded in a 3D tensor with 3 channels corresponding to its RGB values.

2.2.2 Dataset loading and batch splitting

For both the training and the validation datasets a list is created containing the path of their images. Every element in the lists is then mapped to an image-label pair. The lists are then used to create the actual dataset variables.

Each dataset is divided into batches. After some tweaking, a number of 32 elements per batch proved to be optimal for the considered application.

Each dataset is set to repeat, allowing for multiple epochs of training on the same batches. Moreover, the training dataset is set to be prefetched: this allows the next batches to be loaded while the current one is being used, drastically reducing latency and improving throughput at the cost using additional memory.

Only the training dataset is shuffled as it is not really necessary with the validation dataset, since it is only used for evaluation.

3 Experiments and results

Different models have been tested during the developing process; in this section some of those will be shown and the relative results will be discussed.

3.1 Model summary

For each tested model the following data will be reported:

- The NN architecture
- Hyperparameters used
- Data on accuracy for three repeated runs
- Graph of one of the runs
- Comments on the architecture and results

In the layer tables the input layer will not be reported, as it always corresponds to resized image size (144,256,3).

Also note that some abbreviations are used in the Layer Config field in order for the table to fit:

- `k` stands for `kernel size`
- `s` stands for `strides`
- `f` stands for `filters`
- `p` stands for `pool size`
- `r` stands for `rate`

3.2 Models

3.2.1 Baseline Model

Layer Type	Layer Config	Activation	Output	Params
Convolution(Conv2d)	k=5, s=3, f=5	relu	48,86,5	380
Flatten(Flatten)	/	relu	20640	0
Dense(Dense)	u=64	relu	64	1321024
Dense(Dense)	u=6	softmax	64	390
TOTAL				1,321,794

Param	Value
Batch Size	32
Optimizer	Adam
Base lr	0.001
Epochs	20




Figure 2: Graph of the first run

Run	Loss	V.Loss	Acc.	V.Acc.	Δ Acc.
1	0.0016	0.0530	1.0000	0.9754	0.0246
2	0.0012	0.0342	1.0000	0.9821	0.0179
3	0.0060	0.0497	0.9996	0.9866	0.0130
Avg	0.0029	0.0456	0.9996	0.9814	0.0185

This is a very bare-bone model, used to check that everything is working as intended in the network and to get a reference for the next iterations. Nevertheless, we can make a few observations:

- There is some overfitting occurring. This is due to the fact that the image is only slightly reduced in the convolution, thereby creating the need for a substantial number of parameters between the flat layer and the first dense layer.
- The training basically stalls after the 10th epoch as it has already reached maximum accuracy on the training dataset.

3.2.2 Convolution Model

Layer Type	Layer Config	Activation	Output	Params
Convolution(Conv2d)	k=5, s=3, f=5	relu	48,86,5	380
Convolution(Conv2d)	k=5, s=2, f=8	relu	24,43,8	1008
Convolution(Conv2d)	k=3, s=1, f=12	relu	24,43,12	876
Convolution(Conv2d)	k=3, s=1, f=15	relu	24,43,15	1635
Convolution(Conv2d)	k=3, s=1, f=18	relu	24,43,18	2448
Flatten(Flatten)	/	/	20640	0
Dense(Dense)	u=64	relu	64	1188928
Dense(Dense)	u=6	softmax	6	390
TOTAL				1,195,665

Param	Value
Batch Size	32
Optimizer	Adam
Base lr	0.001
Epochs	20




Figure 3: Graph of the third run

Run	Loss	V.Loss	Acc.	V.Acc.	Δ Acc.
1	1.0908e-04	0.0439	1.0000	0.9866	0.0134
2	8.7761e-05	0.0912	1.0000	0.9888	0.0112
3	1.8083e-04	0.0609	1.0000	0.9911	0.0089
Avg	1.2589e-04	0.0653	1.0000	0.9889	0.0112

This model is an iteration of the baseline model, achieved by adding four more convolutions. The convolutions have been added with the following logic in mind:

- The first two convolutions have a stride greater than 1 and that leads to a minimal reduction in the size of the tensor. This in turn should limit the number of parameters needed for training and in turn reduce overfitting.
- The layers are set in a way to progressively increase the number of channels and reduce height and width.

The results show the model is slightly less overfitting than the baseline. Both models reached 1.0000 accuracy on the training set and this leads to an increased validation accuracy.

3.2.3 Convolution and Pooling Model

Layer Type	Layer Config	Activation	Output	Params
Convolution(Conv2d)	k=5, s=1, f=5	relu	144,256,5	380
MaxPooling(MaxPooling2D)	p=2x2	/	72,128,8	0
Convolution(Conv2d)	k=5, s=1, f=8	relu	72,128,8	1008
MaxPooling(MaxPooling2D)	p=2x2	/	36,64,12	0
Convolution(Conv2d)	k=3, s=1, f=12	relu	36,64,12	876
MaxPooling(MaxPooling2D)	p=2x2	/	18,32,15	0
Convolution(Conv2d)	k=3, s=1, f=15	relu	18,32,15	1635
MaxPooling(MaxPooling2D)	p=2x2	/	9,16,18	0
Convolution(Conv2d)	k=3, s=1, f=18	relu	9,16,18	2448
Flatten(Flatten)	/	/	2592	0
Dense(Dense)	u=64	relu	64	165952
Dense(Dense)	u=6	softmax	6	390
TOTAL				172,689

Param	Value
Batch Size	32
Optimizer	Adam
Base lr	0.001
Epochs	20




Figure 4: Graph of the first run

Run	Loss	V.Loss	Acc.	V.Acc.	Δ Acc.
1	1.8558e-04	0.0429	1.0000	0.9933	0.0067
2	6.4439e-04	0.0461	1.0000	0.9866	0.0134
3	1.2224e-04	0.0133	1.0000	0.9933	0.0067
Avg	3.1740e-04	0.0341	1.0000	0.9911	0.0893

This model applies the use of the Conv-Pool technique in order to convolve and reduce the data at every step. Since the reduction is performed by various MaxPooling layers, the first two convolutions have now a stride equal to 1. Pooling helps by reducing the tensor dimension while retaining most of the information. Moreover, the model has a number of parameters which are an order of magnitude lower than the previous model. Therefore the model has now little to no overfitting and has gained a good amount of validation accuracy.

3.2.4 Convolution and Pooling Model with Dropout

Layer Type	Layer Config	Activation	Output	Params
Convolution(Conv2d)	k=5, s=1, f=5	relu	144,256,5	380
MaxPooling(MaxPooling2D)	p=2x2	/	72,128,8	0
Convolution(Conv2d)	k=5, s=1, f=8	relu	72,128,8	1008
MaxPooling(MaxPooling2D)	p=2x2	/	36,64,12	0
Convolution(Conv2d)	k=3, s=1, f=12	relu	36,64,12	876
MaxPooling(MaxPooling2D)	p=2x2	/	18,32,15	0
Convolution(Conv2d)	k=3, s=1, f=15	relu	18,32,15	1635
MaxPooling(MaxPooling2D)	p=2x2	/	9,16,18	0
Convolution(Conv2d)	k=3, s=1, f=18	relu	9,16,18	2448
Dropout(Dropout)	r=0.75	/	9,16,18	0
Flatten(Flatten)	/	/	2592	0
Dense(Dense)	u=64	relu	64	165952
Dropout(Dropout)	r=0.75	/	64	0
Dense(Dense)	u=6	softmax	6	390

Param	Value
Batch Size	32
Optimizer	Adam
Base lr	0.001
Epochs	20




Figure 5: Graph of the first run

Run	Loss	V.Loss	Acc.	V.Acc.	Δ Acc.
1	0.0582	0.0666	0.9817	0.9821	-0.0004
2	0.0063	0.0269	0.9987	0.9933	0.0054
3	0.1172	0.0615	0.9654	0.9866	-0.0212
Avg	0.0606	0.0517	0.9820	0.9873	-0.0054

The next iteration of the model added dropout regularization after the two layers with the most trainable parameters. The desired goal was to chip away another bit off the difference between the training set and the validation set. While this definitely worked, as seen by the negative average Δ accuracy, it came with a slight reduction in the overall validation accuracy.

3.2.5 Batch Normalization Model

Layer Type	Layer Config	Activation	Output	Params
Convolution(Conv2d)	k=5, s=1, f=5	/	144,256,5	375
MaxPooling(MaxPooling2D)	p=2x2	/	72,128,8	0
Batch Norm.(BatchN.)	/	/	72,128,8	15
Relu Activation				
Convolution(Conv2d)	k=5, s=1, f=8	/	72,128,8	1000
MaxPooling(MaxPooling2D)	p=2x2	/	36,64,12	0
Batch Norm.(BatchN.)	/	/	36,64,12	24
Relu Activation				
Convolution(Conv2d)	k=3, s=1, f=12	/	36,64,12	864
MaxPooling(MaxPooling2D)	p=2x2	/	18,32,15	0
Batch Norm.(BatchN.)	/	/	18,32,15	36
Relu Activation				
Convolution(Conv2d)	k=3, s=1, f=15	/	18,32,15	1620
MaxPooling(MaxPooling2D)	p=2x2	/	9,16,18	0
Batch Norm.(BatchN.)	/	/	9,16,18	45
Relu Activation				
Dropout(Dropout)	r=0.06	/	9,16,18	0
Convolution(Conv2d)	k=3, s=1, f=18	/	9,16,18	2430
Flatten(Flatten)	/	/	2592	0
Batch Norm.(BatchN.)	/	/	2592	7776
Relu Activation				
Dense(Dense)	u=64	/	64	165888
Batch Norm.(BatchN.)	/	/	64	192
Relu Activation				
Dropout(Dropout)	r=0.06	/	64	0
Dense(Dense)	u=6	softmax	6	390
TOTAL				180,655

Param	Value
Batch Size	32
Optimizer	Adam
Base lr	0.00005
Epochs	50

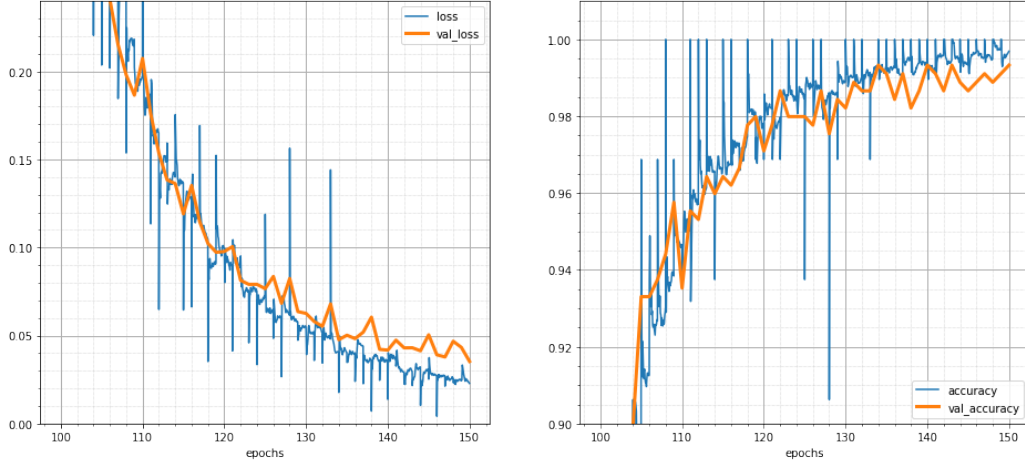


Figure 6: Graph of the third run

Run	Loss	V.Loss	Acc.	V.Acc.	Δ Acc.
1	0.0276	0.0439	0.9962	0.9933	0.0029
2	0.0255	0.0510	0.9962	0.9888	0.0074
3	0.0227	0.0353	0.9969	0.9933	0.0036
Avg	0.0253	0.0434	0.9964	0.9918	0.0046

This iteration of the model adds batch normalization in order to further improve general performance.

Adding batch normalization caused the model to train much faster, but this presented a problem: the training was very biased towards the training set from the first epoch. This led to a very big gap in validation and train accuracy, in the order of 0.75 against 0.14 after the first epoch.

To counter this negative effect, the starting learning rate was reduced to 0.00005, which allowed for much better results. The slower learning rate called for an increased number of epochs.

3.3 Models comparison

When deciding which is best, the following factors should be taken into account:

- **Accuracy:** the more accurate the model, the better. The measure to look for is validation accuracy.

- **Generalization:** the model needs to work with images outside of the training dataset. If the model generalizes well, the difference between its accuracy and validation accuracy should be little to none.
- **Number of parameters:** if two models have similar performances, the one with less parameters is to be preferred. Having less parameters generally makes the model faster in training and evaluation. It also comes with the added benefit of having a lower chance of overfitting.

#	Loss	V.Loss	Acc.	V.Acc	Δ Acc.	Params.
1	0.0029	0.0456	0.9996	0.9814	0.0185	1,321,794
2	0.0001	0.0653	1.0000	0.9889	0.0112	1,195,665
3	0.0003	0.0341	1.0000	0.9911	0.0893	172,689
4	0.0606	0.0517	0.9820	0.9873	-0.0054	172,689
5	0.0253	0.0434	0.9964	0.9918	0.0046	180,655

Table 1: Average performance of the various model

On this grounds it's clear that models 4 and 5 are the best candidates for the given dataset.

4 Conclusions

As seen in the previous sections, after a few experiments, this project has been succesful in classifying the images of Turkish Lira banknotes with high accuracy and low data loss. (altro?) aggiungere migliori risultati ottenuti -
Esempio di citazione[?]

References

- [1] Easiest way to download kaggle data in Google Colab
<https://www.kaggle.com/general/74235>
- [2] Load images
https://www.tensorflow.org/tutorials/load_data/images
- [3] Classification
<https://www.tensorflow.org/tutorials/images/classification>
- [4] Convolutional Neural Network (CNN)
<https://www.tensorflow.org/tutorials/images/cnn>
- [5] A Comprehensive Guide to Convolutional Neural Networks — the ELI5 way
<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>
- [6] A guide to an efficient way to build neural network architectures- Part II: Hyper-parameter selection and tuning for Convolutional Neural Networks using Hyperas on Fashion-MNIST
<https://towardsdatascience.com/a-guide-to-an-efficient-way-to-build-neural-network-architectures-part-ii-hyper-parameter-42efca01e5d73>
- [7] Prof. Notebook on deep learning