

Documentation for Algorithms for Massive Datasets project: Turkish lira recognizer

Caccaro Sebastiano
Cavagnino Matteo
A.A.2019/2020

We declare that this material, which We now submit for assessment, is entirely our own work and has not been taken from the work of others, save and to the extent that such work has been cited and acknowledged within the text of our work. We understand that plagiarism, collusion, and copying are grave and serious offences in the university and accept the penalties that would be imposed should we engage in plagiarism, collusion or copying. This assignment, or any part of it, has not been previously submitted by us or any other person for assessment on this or any other course of study.

Contents

1	Introduction	1
2	The Turkish Lira banknotes dataset	2
2.1	Preprocessing techniques applied to the dataset	3
2.2	Considered algorithms and their implementation	4
3	Scalability of the proposed solution	5
4	Experiments and results	6
4.1	Model summary	6
4.2	Models	6
4.2.1	Baseline Model	6
4.2.2	Convolution Model	8
4.2.3	Convolution and Pooling Model	9
4.2.4	Convolution and Pooling Model with Dropout	11
4.2.5	Batch Normalization Model	12
5	Conclusions	14
	References	15

1 Introduction

The Objective of this project is to build a Turkish Lira banknotes image recognizer through a Convolutional Neural Network. The proposed solution, based on Tensorflow libraries, contains steps to dynamically download the dataset, preprocess the images in it and use the processed images to train a Convolutional Neural Network in recognizing and classifying them. Since the given dataset contains many images and since the request is to classify some precise details of them, it's expected to achieve good results from the proposed solution and in particular from the proposed model.

2 The Turkish Lira banknotes dataset

The chosen dataset <add ref> is originally composed of 6000 images of Turkish Lira banknotes, organized in folders grouping banknotes by their value and already splitted in training and validation set.

2.1 Preprocessing techniques applied to the dataset

In this section the following preprocessing techniques applied to the considered dataset will be discussed:

- image scaling
- training and validation sets creation

Regarding the image scaling, the given images have a size of 720x1280; to not overload the memory of the computing machine it has been reduced by 5 times resulting in a size of 140x256 per image. For the training and validation sets, they have been created starting from two text files, provided with the dataset, listing all the images that needed to be used for the training or validation phases. The training dataset array has then been processed using shuffling, batching and repeating techniques to prepare it for the learning process. The validation dataset array has instead been processed using only batching and repeating; for both the datasets, the batches size has been set to 32 images per batch.

The training dataset has also been processed using the prefetch technique, this allows later elements to be prepared while the current element is being processed. This often improves latency and throughput, at the cost of using additional memory to store prefetched elements.

2.2 Considered algorithms and their implmentation

?

3 Scalability of the proposed solution

The Scalability of this project is granted by (batching, caching, img scaling,
? ...)

4 Experiments and results

Different models have been tested during the developing process; in this section some of those will be shown and the relative results will be discussed.

4.1 Model summary

For each tested model the following data will be reported:

- The NN architecture
- Hyperparameters used
- Data on accuracy for three repeated runs
- Graph of one of the runs
- Comment on the architecture and results

In the layer tables the input layer will not be reported, as it always corresponds to resized image size (144,256,3).

Also note that some abbreviations are used in the Layer Config field in order to for the table to fit:

- `k` stands for `kernel size`
- `s` stands for `strides`
- `f` stands for `filters`
- `p` stands for `pool size`
- `r` stands for `rate`

4.2 Models

4.2.1 Baseline Model

Layer Type	Layer Config	Activation	Output	Params
Convolution(Conv2d)	k=5, s=3, f=5	relu	48,86,5	380
Flatten(Flatten)	/	relu	20640	0
Dense(Dense)	u=64	relu	64	1321024
Dense(Dense)	u=6	softmax	64	390
			TOTAL	1,321,794

Param	Value
Batch Size	32
Optimizer	Adam
Base lr	0.001
Epochs	20




Figure 1: Graph of the first run

Run	Loss	V.Loss	Acc.	V.Acc.	Δ Acc.
1	0.0016	0.0530	1.0000	0.9754	0.0246
2	0.0012	0.0342	1.0000	0.9821	0.0179
3	0.0060	0.0497	0.9996	0.9866	0.0130
Avg	0.0029	0.0456	0.9996	0.9814	0.0185

This is a very bare-bone, used to check everything is working in the net and to get a reference for the next iterations. Nevertheless, we can make a few observation:

- There is some overfitting occurring. This is due to the fact the image is only slightly reduced in the convolution, thereby creating the need for a substantial number of parameters between the flat layer and the first dense layer.
- The training basically stall after the 10th epoch, as it has already reach maximum accuracy on the train dataset.

4.2.2 Convolution Model

Layer Type	Layer Config	Activation	Output	Params
Convolution(Conv2d)	k=5, s=3, f=5	relu	48,86,5	380
Convolution(Conv2d)	k=5, s=2, f=8	relu	24,43,8	1008
Convolution(Conv2d)	k=3, s=1, f=12	relu	24,43,12	876
Convolution(Conv2d)	k=3, s=1, f=15	relu	24,43,15	1635
Convolution(Conv2d)	k=3, s=1, f=18	relu	24,43,18	2448
Flatten(Flatten)	/	/	20640	0
Dense(Dense)	u=64	relu	64	1188928
Dense(Dense)	u=6	softmax	6	390
TOTAL				1,195,665

Param	Value
Batch Size	32
Optimizer	Adam
Base lr	0.001
Epochs	20




Figure 2: Graph of the third run

Run	Loss	V.Loss	Acc.	V.Acc.	Δ Acc.
1	1.0908e-04	0.0439	1.0000	0.9866	0.0134
2	8.7761e-05	0.0912	1.0000	0.9888	0.0112
3	1.8083e-04	0.0609	1.0000	0.9911	0.0089
Avg	1.2589e-04	0.0653	1.0000	0.9889	0.0112

This models is an iteration of the baseline model, achieved by adding four more convolution. The convolution have been added with the following logic in mind:

- The first two convolutions have a stride greater than 1, and that leads to a minimal reductions in the size of the tensor. This in turn should limit the number of parameters needed for training and in turn reduce overfitting.
- The layers are set in way to progressively increase the number of channel and reduce height and width.

The results show the model is slightly less overfitting than the baseline. Both models reached 1.0000 accuracy on the training set, so this translate in an increased validation accuracy.

4.2.3 Convolution and Pooling Model

Layer Type	Layer Config	Activation	Output	Params
Convolution(Conv2d)	k=5, s=1, f=5	relu	144,256,5	380
MaxPooling(MaxPooling2D)	p=2x2	/	72,128,8	0
Convolution(Conv2d)	k=5, s=1, f=8	relu	72,128,8	1008
MaxPooling(MaxPooling2D)	p=2x2	/	36,64,12	0
Convolution(Conv2d)	k=3, s=1, f=12	relu	36,64,12	876
MaxPooling(MaxPooling2D)	p=2x2	/	18,32,15	0
Convolution(Conv2d)	k=3, s=1, f=15	relu	18,32,15	1635
MaxPooling(MaxPooling2D)	p=2x2	/	9,16,18	0
Convolution(Conv2d)	k=3, s=1, f=18	relu	9,16,18	2448
Flatten(Flatten)	/	/	2592	0
Dense(Dense)	u=64	relu	64	165952
Dense(Dense)	u=6	softmax	6	390
			TOTAL	172,689

Param	Value
Batch Size	32
Optimizer	Adam
Base lr	0.001
Epochs	20




Figure 3: Graph of the first run

Run	Loss	V.Loss	Acc.	V.Acc.	Δ Acc.
1	1.8558e-04	0.0429	1.0000	0.9933	0.0067
2	6.4439e-04	0.0461	1.0000	0.9866	0.0134
3	1.2224e-04	0.0133	1.0000	0.9933	0.0067
Avg	3.1740e-04	0.0341	1.0000	0.9911	0.0893

This model applies the use of the Conv-Pool technique in order to convolve and reduce the data at every step. Since the reduction is performed by various MaxPooling layers, the first two convolution have now a stride equal to 1. Pooling helps by reducing the tensor dimension while retaining most of the information. Moreover, the model has a number of parameter which an order of magnitude lower than the previous model. Therefore the model has now little to no overfitting, and has gained a good amount of validation accuracy.

4.2.4 Convolution and Pooling Model with Dropout

Layer Type	Layer Config	Activation	Output	Params
Convolution(Conv2d)	k=5, s=1, f=5	relu	144,256,5	380
MaxPooling(MaxPooling2D)	p=2x2	/	72,128,8	0
Convolution(Conv2d)	k=5, s=1, f=8	relu	72,128,8	1008
MaxPooling(MaxPooling2D)	p=2x2	/	36,64,12	0
Convolution(Conv2d)	k=3, s=1, f=12	relu	36,64,12	876
MaxPooling(MaxPooling2D)	p=2x2	/	18,32,15	0
Convolution(Conv2d)	k=3, s=1, f=15	relu	18,32,15	1635
MaxPooling(MaxPooling2D)	p=2x2	/	9,16,18	0
Convolution(Conv2d)	k=3, s=1, f=18	relu	9,16,18	2448
Dropout(Dropout)	r=0.75	/	9,16,18	0
Flatten(Flatten)	/	/	2592	0
Dense(Dense)	u=64	relu	64	165952
Dropout(Dropout)	r=0.75	/	64	0
Dense(Dense)	u=6	softmax	6	390

Param	Value
Batch Size	32
Optimizer	Adam
Base lr	0.001
Epochs	20




Figure 4: Graph of the first run

Run	Loss	V.Loss	Acc.	V.Acc.	Δ Acc.
1	0.0582	0.0666	0.9817	0.9821	-0.0004
2	0.0063	0.0269	0.9987	0.9933	0.0054
3	0.1172	0.0615	0.9654	0.9866	-0.0212
Avg	0.0606	0.0517	0.9820	0.9873	-0.0054

4.2.5 Batch Normalization Model

Layer Type	Layer Config	Activation	Output	Params
Convolution(Conv2d)	k=5, s=1, f=5	/	144,256,5	375
MaxPooling(MaxPooling2D)	p=2x2	/	72,128,8	0
Batch Norm.(BatchN.)	/	/	72,128,8	15
Relu Activation				
Convolution(Conv2d)	k=5, s=1, f=8	/	72,128,8	1000
MaxPooling(MaxPooling2D)	p=2x2	/	36,64,12	0
Batch Norm.(BatchN.)	/	/	36,64,12	24
Relu Activation				
Convolution(Conv2d)	k=3, s=1, f=12	/	36,64,12	864
MaxPooling(MaxPooling2D)	p=2x2	/	18,32,15	0
Batch Norm.(BatchN.)	/	/	18,32,15	36
Relu Activation				
Convolution(Conv2d)	k=3, s=1, f=15	/	18,32,15	1620
MaxPooling(MaxPooling2D)	p=2x2	/	9,16,18	0
Batch Norm.(BatchN.)	/	/	9,16,18	45
Relu Activation				
Dropout(Dropout)	r=0.06	/	9,16,18	0
Convolution(Conv2d)	k=3, s=1, f=18	/	9,16,18	2430
Flatten(Flatten)	/	/	2592	0
Batch Norm.(BatchN.)	/	/	2592	7776
Relu Activation				
Dense(Dense)	u=64	/	64	165888
Batch Norm.(BatchN.)	/	/	64	192
Relu Activation				
Dropout(Dropout)	r=0.06	/	64	0
Dense(Dense)	u=6	softmax	6	390
			TOTAL	180,655

Param	Value
Batch Size	32
Optimizer	Adam
Base lr	0.00005
Epochs	50




Figure 5: Graph of the third run

Run	Loss	V.Loss	Acc.	V.Acc.	Δ Acc.
1	0.0276	0.0439	0.9962	0.9933	0.0029
2	0.0255	0.0510	0.9962	0.9888	0.0074
3	0.0227	0.0353	0.9969	0.9933	0.0036
Avg	0.0253	0.0434	0.9964	0.9918	0.0046

5 Conclusions

As seen in the previous sections, after a few experiments, this project has been succesful in classifying the images of Turkish Lira banknotes with high accuracy and low data loss. (altro?) aggiungere migliori risultati ottenuti -
Esempio di citazione[1]

References

- [1] Michel Goossens, Frank Mittelbach, and Alexander Samarin. *The L^AT_EX Companion*. Addison-Wesley, Reading, Massachusetts, 1993.
- [2] Albert Einstein. *Zur Elektrodynamik bewegter Körper*. (German) [*On the electrodynamics of moving bodies*]. Annalen der Physik, 322(10):891–921, 1905.
- [3] Knuth: Computers and Typesetting,
<http://www-cs-faculty.stanford.edu/~uno/abcde.html>